

```

/*****
Authors : Albert Owusu-Asare , Box 4497, <owusuasa@grinnell.edu>
         : Ezra Edgerton , Box 3503, <edgerton@grinnell.edu>

This document contains answers to questions from CS213 Lab: Threads

Questions:
http://www.cs.grinnell.edu/~weinman/courses/CSC213/2014F/labs/threads.html

modified code has the MODIFIED keyword
Original author is cited below :
* Jerod Weinman
* 21 May 2008
*****/

#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#include "matrix.h"
#include "matrixop.h"

#define SUB2IND(row,col,numcols) (row) * (numcols) + (col)

/* Check the dimensions of three matrices for the sum C=A+B
*
* Returns -1 if there is a dimension mismatch, and zero otherwise. */
int mtxCheckAddDim(const struct matrix_t *a, const struct matrix_t *b,
                  const struct matrix_t *c)
{
    if (a->rows != b->rows || a->rows != c->rows ||
        a->cols != b->cols || a->cols != c->cols ) {
        return -1;
    }

    return 0;
}

/* Matrix addition C = A + B
*
* Preconditions:
*   Matrix parameters a,b and c all have the same dimension
*
* Postconditions:
*   The resulting matrix sum is stored in parameter c
*   Return value of 0 indicates successful completion of the addition
*/
int mtxAdd( const struct matrix_t *a, const struct matrix_t *b,
            struct matrix_t *c) {
    /* Make sure these matrices are "addable" */
    int res = mtxCheckAddDim(a,b,c);

    if (res<0) {
        fprintf(stderr,"Dimension mismatch for matrix add.\n");
        return res;
    }

    int m = a->rows;          /* Dimensions of input matrices */
    int n = a->cols;
    MTX_TYPE *pa, *pb, *pc; /* Pointers to matrix data */

```

```

    pa = a->data;              /* Initialize pointers to beginning of data */
    pb = b->data;
    pc = c->data;

    /* Iterate over every entry in the matrices, adding and storing the result */
    int i,j;
    for (i=0 ; i<m ; i++)
        for (j=0 ; j<n ; j++, pa++,pb++,pc++)
            *pc = *pa + *pb;

    return 0; /* Indicate successful completion */
}

/* MODIFIED : This is an additional piece to the original
* Check the dimensions of three matrices for the sum C=AB
*
* Returns -1 if there is a dimension mismatch, and zero otherwise. */
int mtxCheckMultiplyDim(const struct matrix_t *a, const struct matrix_t *b,
                       const struct matrix_t *c)
{
    if (a->cols != b->rows || a->rows != c->rows
        || b->cols != c->cols ) {
        return -1;
    }

    return 0;
}

/*
* MODIFIED
* Matrix multiplication C = AB
*
* Preconditions:
*   Matrix parameters a,b,c have the required dimension.
*   mtxCheckMultiplyDim(a,b,c) returns >0
*
* Postconditions:
*   The resulting matrix sum is stored in parameter c
*   Return value of 0 indicates successful completion of the multiplication
*/
int mtxMultiplyMin( const struct matrix_t *a, const struct matrix_t *b,
                   struct matrix_t *c)
{
    /* Make sure these matrices are "multiply-able" */
    int res = mtxCheckMultiplyDim(a,b,c);

    if (res<0) {
        fprintf(stderr,"Dimension mismatch for matrix add.\n");
        return res;
    }

    int m = a->rows;          /* Dimensions of input matrices */
    int n = a->cols;
    int p = b->cols;

    MTX_TYPE *pa, *pb, *pc; /* Pointers to matrix data */

    pa = a->data;              /* Initialize pointers to beginning of data */
    pb = b->data;

```

```

pc = c->data;

/* Iterate over every entry in the matrices, adding and storing the result */
int i,j,k, l;

/* initialise elements of matrix c to 0*/

for (l = 0; l < (m*p); l++, pc++)
{
    *pc = 0;
}
pc = c->data; // point back to start of c matrix

for (i = 0; i < m; i++, pc += SUB2IND(1,0,p), pb = b->data)
//i loops through the rows of a and the rows of c
{
    for (k = 0; k < n; k++, pa++, pc = pc - p)
        //k loops through the columns of a and the rows of b
        {
            for (j = 0; j < p; j++,pb++,pc++)
                //the columns of b and the columns of c
                {
                    *pc += *pa * *pb;
                }
            }
        }

    }

return 0; /* Indicate successful completion */
}
/*
 * MODIFIED
 * Does the matrix multiplication work for one thread
 * preconditions: none
 * postconditions : matrix c is modified such that each of its values agrees with m
atrix
 * multiplication of c = a * b.
 *
 * much of the code involving the creation and running of threads was taken from th
e POSIX
 * Threads Programming tutorial:
 * https://chttps://computing.llnl.gov/tutorials/pthreads/#Joining
 * computing.llnl.gov/tutorials/pthreads/#PassingArguments
 */

void* threadMtxMultiplyMin(void *multParam)
{
    struct matrixThreadParam_t *data;
    data = ( struct matrixThreadParam_t *) multParam ;

    /* Dimensions of input matrices */
    int m = data->a->rows;
    int n = data->a->cols;
    int p = data->b->cols;

    int stride = data->numThreads;
    int ID = data->threadId;

    /* Pointers to matrix data */
    MTX_TYPE *pa, *pb, *pc;

    /* Initialize pointers to beginning of data */
    pa = data->a->data;
    pb = data->b->data;
    pc = data->c->data;

```

```

/* Iterate over every entry in the matrices, adding and storing the result */
int i,j,k;

pc += SUB2IND(ID, 0, p);
pa += SUB2IND(ID, 0, n);

for (i = ID; i < m;
     i= i + stride, pc += SUB2IND(stride, 0, p), pb = data->b->data,
     pa += SUB2IND(stride-1, 0 , n))
//i loops through the rows of a and the rows of c
{
    for (k = 0; k < n; k++, pa++, pc = pc - p)
        //k loops through the columns of a and the rows of b
        {
            for (j = 0; j < p; j++,pb++,pc++)
                //the columns of b and the columns of c
                {
                    *pc += *pa * *pb;
                }
            }
        }

    return NULL; /* Indicate successful completion */
}

/*
 * MODIFIED
 * Computes the threaded concurrent matrix product.
 * preconditions: none
 * postconditions : matrix c is modified such that each of its values agrees with m
atrix
 * multiplication of c = a * b.
 *
 * much of the code involving the creation and running of threads was taken from th
e POSIX
 * Threads Programming tutorial:
 * https://chttps://computing.llnl.gov/tutorials/pthreads/#Joining
 * computing.llnl.gov/tutorials/pthreads/#PassingArguments
 */

int parMtxMultiply( const struct matrix_t *a,
                    const struct matrix_t *b,
                    struct matrix_t *c,
                    int numThreads)
{
    //checks to see if matrixes are multipliable
    int res = mtxCheckMultiplyDim(a,b,c);

    if (res<0) {
        fprintf(stderr,"Dimension mismatch for matrix add.\n");
        return res;
    }

    /* initialise elements of matrix c to 0*/
    int l;
    MTX_TYPE *pc;
    pc = c->data;

    for (l = 0; l < (a->rows * b->cols); l++, pc++)
    {
        *pc = 0;
    }
    pc = c->data; // point back to start of c matrix

```

```
//creates threads according to numThreads
int t;
int rc; //gets return value when thread is created and when thread is joined
struct matrixThreadParam_t thread_matrix_array[numThreads];
pthread_t threads[numThreads];

pthread_attr_t attr;
void *status;

//initialization of attr to make threads joinable
pthread_attr_init(&attr);
pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);

for (t = 0; t < numThreads; t++)
{
    //set elements of thread_array[t]
    thread_matrix_array[t].a = a;
    thread_matrix_array[t].b = b;
    thread_matrix_array[t].c = c;
    thread_matrix_array[t].numThreads = numThreads;
    thread_matrix_array[t].threadId = t;
    rc = pthread_create(&threads[t], &attr, threadMtxMultiplyMin,
                       (void *) &thread_matrix_array[t]);

    if (rc) {
        fprintf(stderr, "Could not create thread %d\n", t);
        return EXIT_FAILURE;
    }
}
//join threads
pthread_attr_destroy(&attr);
for (t = 0; t < numThreads; t++)
{
    rc = pthread_join(threads[t], &status);
    if (rc)
    {
        fprintf(stderr, "Thread %d could not join\n", t);
        return EXIT_FAILURE;
    }
}
return 0;
}
```

```

/*****
Authors : Albert Owusu-Asare , Box 4497, <owusuasa@grinnell.edu>
        : Ezra Edgerton , Box 3503, <edgerton@grinnell.edu>
*****/

```

This document contains answers to questions from CS213 Lab: Threads

Questions:
<http://www.cs.grinnell.edu/~weinman/courses/CSC213/2014F/labs/threads.html>

modified code has the MODIFIED keyword

* Created by Jerod Weinman, 21 May 2008

```

*****/

```

```

#ifdef __MATRIXOP_H__
#define __MATRIXOP_H__

```

```

#include "matrix.h"

```

```

/* Matrix addition C = A + B

```

```

*
* Preconditions:
*   Matrix parameters a,b and c all have the same dimension
*
* Postconditions:
*   The resulting matrix sum is stored in parameter c
*   Return value of 0 indicates successful completion of the addition
*/

```

```

int mtxAdd( const struct matrix_t *a, const struct matrix_t *b,
            struct matrix_t *c );

```

```

/*
* MODIFIED
* Matrix multiplication C = AB
*
* Preconditions:
*   Matrix parameters a,b,c have the required dimension.
*   mtxCheckMultiplyDim(a,b,c) returns >0
*
* Postconditions:
*   The resulting matrix sum is stored in parameter c
*   Return value of 0 indicates successful completion of the multiplication
*/

```

```

int mtxMultiplyMax( const struct matrix_t *a, const struct matrix_t *b,
                   struct matrix_t *c );

```

```

int mtxMultiplyMin( const struct matrix_t *a, const struct matrix_t *b,
                   struct matrix_t *c );

```

```

/*
*MODIFIED
* struct code taken from Jerod Weinman's CSC213 Threads lab:
* http://www.cs.grinnell.edu/~weinman/courses/CSC213/2014F/labs/threads.html

```

```

* matrixThreadParam_t contains the parameters needed for a threaded matrix
* product.
*/

```

```

struct matrixThreadParam_t {
    const struct matrix_t *a;
    const struct matrix_t *b;

```

```

    struct matrix_t *c;
    int numThreads;
    int threadId;
};

```

```

/*
* MODIFIED
* Computes the threaded concurrent matrix product.
* preconditions: none
* postconditions : matrix c is modified such that each of its values agrees with m
atrix
* multiplication of c = a * b.
*
* much of the code involving the creation and running of threads was taken from th
e POSIX
* Threads Programming tutorial:
* https://chttps://computing.llnl.gov/tutorials/pthreads/#Joining
* computing.llnl.gov/tutorials/pthreads/#PassingArguments
*/

```

```

int parMtxMultiply( const struct matrix_t *a,
                   const struct matrix_t *b,
                   struct matrix_t *c,
                   int numThreads );

```

```

/*
* MODIFIED
* Does the matrix multiplication work for one thread
* preconditions: none
* postconditions : matrix c is modified such that each of its values agrees with m
atrix
* multiplication of c = a * b.
*
* much of the code involving the creation and running of threads was taken from th
e POSIX
* Threads Programming tutorial:
* https://chttps://computing.llnl.gov/tutorials/pthreads/#Joining
* computing.llnl.gov/tutorials/pthreads/#PassingArguments
*/

```

```

void* threadMtxMultiplyMin(void *multParam);

```

```

#endif

```

Script started on Tue 14 Oct 2014 06:45:13 PM CDT

burroughs\$ make

gcc -g -Wall -D_GNU_SOURCE -c matrixop.c

burroughs\$ make pt\007imemultiply

gcc -g -Wall -D_GNU_SOURCE -c matrixop.c

gcc -g -Wall -D_GNU_SOURCE -c matrixutil.c

gcc -g -Wall -D_GNU_SOURCE -pthread -o ptimemultiply \

matrixop.o matrixutil.o ptimemultiply.c

burroughs\$ p 033[K 033[K 007 007 007 007exit

Script done on Tue 14 Oct 2014 06:45:28 PM CDT

Data from running :

```
for n in 128 192 256 384 512 768 1024 1536 2048;
do echo -ne $nt >> mtxMultiplyProfile;
./ptimemultiply $n 1 >> mtxMultiplyProfile; done

for n in 128 192 256 384 512 768 1024 1536 2048;
do echo -ne $nt >> mtxMultiplyProfile;
./ptimemultiply $n 2 >> mtxMultiplyProfile; done

for n in 128 192 256 384 512 768 1024 1536 2048;
do echo -ne $nt >> mtxMultiplyProfile;
./ptimemultiply $n 3 >> mtxMultiplyProfile; done

for n in 128 192 256 384 512 768 1024 1536 2048;
do echo -ne $nt >> mtxMultiplyProfile;
./ptimemultiply $n 4 >> mtxMultiplyProfile; done
```

Threads: 1

0.010495	0.005843	0.006341
0.028220	0.022967	0.024868
0.053346	0.056714	0.050666
0.161465	0.152446	0.159345
0.366919	0.357090	0.368407
1.204320	1.196387	1.202859
2.934068	2.922738	2.968047
9.922737	9.914455	10.058668
23.468226	23.461429	23.460662

Threads: 2

0.007192	0.007133	0.006954
0.019720	0.015568	0.013420
0.035314	0.031007	0.027983
0.087355	0.085374	0.089545
0.190058	0.197384	0.198516
0.625959	0.629656	0.628030
1.517283	1.515384	1.538852
5.111999	5.110251	5.184066
12.090712	12.086609	12.089291


Threads: 3

0.005107	0.004952	0.005003
0.015987	0.016010	0.013790
0.016739	0.016980	0.021162
0.064651	0.061131	0.060980
0.136234	0.134229	0.134733
0.428568	0.434557	0.432197
1.040933	1.043946	1.053486
3.497636	3.498054	3.537673
8.278215	8.275488	8.278231


Thread: 4

0.004011	0.004247	0.003847
0.008683	0.005844	0.005583
0.012906	0.012855	0.012919
0.042088	0.042247	0.042629
0.098370	0.097934	0.099495
0.329105	0.329907	0.330154
0.798945	0.795852	0.808042
2.688160	2.687280	2.725116
6.355615	6.373270	6.355293

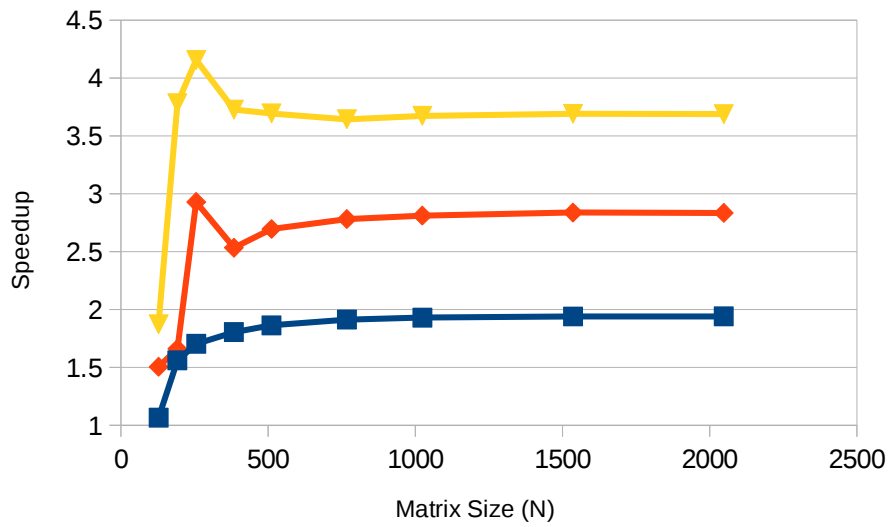
Graph of Average Times vs Matrix Size




Graph of Average Times vs Matrix Size on Logarithmic Scale



Matrix Size vs Speedup




Matrix Size N vs Efficiency




Number of Threads vs Speedup

In terms of largest matrix size tested



Number of Threads vs Efficiency

In terms of largest matrix size tested



```
*****
Authors : Albert Owusu-Asare , Box 4497, <owusuasa@grinnell.edu>
        : Ezra Edgerton      , Box 3503, <edgerton@grinnell.edu>
```

This document contains answers to questions from CS213 Lab: Threads

Questions:

<http://www.cs.grinnell.edu/~weinman/courses/CSC213/2014F/labs/threads.html>

```
*****
```

Answers

Relevant information:

Page size: 4096

```
Number of processors available (nproc) = 4;
processor 0:
cpu MHz = 1600.000
cache size = 3072 KB
cpu cores = 4
address sizes : 36 bits physical, 48 bits virtual
```

```
processor 1:
cpu MHz = 1600.000
cache size = 3072 KB
cpu cores = 4
address sizes : 36 bits physical, 48 bits virtual
```

```
processor 2:
cpu MHz = 1600.000
cache size = 3072 KB
cpu cores = 4
address sizes : 36 bits physical, 48 bits virtual
```

```
processor 3:
cpu MHz = 1600.000
cache size = 3072 KB
cpu cores = 4
address sizes : 36 bits physical, 48 bits virtual
```

Part B

3)

a.

The average times taken **for** a set of threaded processes to run a matrix of size N is the average time **for** 1 thread divided by the number of threads in the process.

As a result the plots of the different threaded processes showed less growth in average **time**(seconds) vs matrix size as the number of threads rose.

As we increase the number of threads, we are able to obtain more concurrency because as we saw from nproc, we have 4 available processors on the computer. The more threads we use, the less number of rows the thread is responsible **for** calculating in the matrix and the greater number of rows can be computed at a time; thus we have less overall time calculating

the matrix. In the **case of** 4 threads being used, each thread uses a processor to calculate the rows that it is responsible **for** so the average run time is 1/4 of the one threaded average run time. This holds **for** the 2 and 3 threaded runs being 1/2 and 1/3 respectively of the run time of one thread as well.

b. The trend in speedup as N increases is that the differently threaded speedups hold more or less constant. Two Threads holds at just below 2, Three threads holds at a little more below 3 than 2 threads held below 2, and four threads holds at a little more below 4 than 3 threads held below 3. The difference from the numbers that the speedups were just below is caused by the overhead from creating and running the threads, the overhead will increase when we create more threads, bringing the speedup time below the round number that just running the program with overheadless threads would be. 4 threads is the farthest below 4 because the 4 threads would have the most overhead.

The best answer we could come up with is that the deviations at the beginning must have something to **do** with the page or cache size. That at an N of between 256 and 384, each row would pass the limit of the page size this would cause a larger number of tlb misses and slow down the process significantly. Before this, the tlb misses of each thread would decrease as the number of threads increased, so the speedup time **for** runs with greater threads would increase more quickly.

c. Efficiency seems to remove the advantage that the speedup metric measures by dividing by the number of threads, it seems to measure the differences in overhead, because with one thread, the efficiency is one. As N increases, we see that 2 threads is the most efficient, holding at a little below 1, **while** 3 and 4 threads are less than the smaller number of threads. This is because, as we said in b, the overheads **for** more threads increases as we build more threads.

At smaller values of N, the deviation from the trend is the same as our answer **for** part b. However, we see that 4 threads goes above 1 in efficiency at N = 256. This must be due to a reason similar to what was mentioned in b.

d.

As P increases the speedup increases linearly at a rate of slightly less than 1. As we said previously, the overhead causes the slight decrease in growth below 1. Without the overhead of thread creation and joining, the growth would be at a constant increase of one because the number of threads decreases the run time of the matrix multiply at that rate, as we explained in 3a above. More threads mean more concurrent processing of rows, and as a result a faster overall calculation of the matrix.

e. As P increases the efficiency decreases linearly at a rate of roughly -.025. This measures the amount of time that is spent in overhead **for** each increase in threads. The more threads we have, the more time we spend creating and joining them, and that cancels out a small amount of the speedup effect.